```
import numpy as np
from time import perf_counter
from sklearn.datasets import make_blobs
from sklearn.cluster import KMeans, MiniBatchKMeans
from sklearn.metrics import silhouette_score
np.random.seed(42)



#toy dataset (five clearly-separated blobs)
X, _ = make_blobs(n_samples=5000,
                  centers=[(0,2.3),(-1.5,2.3),(-2.8,1.8),(-2.8,2.8),(-
                  cluster_std=[0.4,0.3,0.1,0.1,0.1],
                  random_state=42)
print(X.shape)
```

```
(5000, 2)
```

```
#regular K-Means
k = 5
t0 = perf_counter()
km = KMeans(n_clusters=k, random_state=42)
labels_km = km.fit_predict(X)
t_km = perf_counter() - t0

print(f"Regular KMeans inertia: {km.inertia_:.2f}")
print(f"Regular KMeans silhouette: {silhouette_score(X, labels_km):.3f
print(f"Fit time: {t_km*1e3:.1f} ms")
```

```
Regular KMeans inertia: 534.68
Regular KMeans silhouette: 0.639
Fit time: 171.0 ms
```

```
#mini batch (batch) K-Means
t0 = perf_counter()
mb = MiniBatchKMeans(n_clusters=k, batch_size=256, random_state=42)
labels_mb = mb.fit_predict(X)
t_mb = perf_counter() - t0

print(f"MiniBatch inertia: {mb.inertia_:.2f}")
print(f"MiniBatch silhouette: {silhouette_score(X, labels_mb):.3f}")
print(f"Fit time: {t_mb*1e3:.1f} ms")
```

```
MiniBatch inertia: 553.16
MiniBatch silhouette: 0.629
Fit time: 41.4 ms
```

```
#elbow & silhouette helper
import matplotlib.pyplot as plt
inertias = []
silhs = []
for k_test in range(2,11):
    km_tmp = KMeans(n_clusters=k_test, random_state=42).fit(X)
    inertias.append(km_tmp.inertia_)
    silhs.append(silhouette_score(X, km_tmp.labels_))
plt.plot(range(2,11), inertias, 'o-')
plt.xlabel('k'); plt.ylabel('Inertia'); plt.title('Elbow check'); plt.

plt.plot(range(2,11), silhs, 'o-')
plt.xlabel('k'); plt.ylabel('Silhouette'); plt.title('Silhouette check
```

Release notes

TPU support GitHub.

- Now your popular Kaggle datasets are cached for quick retrieval.
- Upgraded Colab runtimes to Python 3.11.

Python package upgrades

- bigframes 1.29.0 -> 1.42.0
- TensorFlow 2.17.1 -> 2.18.0
- tensorboard 2.17.1 -> 2.18.0
- keras 3.5.0 -> 3.8.0
- torch 2.5.1 -> 2.6.0
- torchaudio 2.5.1 -> 2.6.0
- torchvision 0.20.1 -> 0.21.0
- fastai 2.7.18 -> 2.7.19
- ipykernel 5.5.6 -> 6.17.1
- google-genai 0.3.0 -> 1.9.0
- google-auth 2.27.0 -> 2.38.0
- Tornado 6.3.3 -> 6.4.2
- jax 0.4.33 -> 0.5.2
- accelerate 1.2.1 -> 1.5.2
- transformers 4.47.1 -> 4.50.3
- openai 1.57.4 -> 1.70.0
- kagglehub 0.3.6 -> 0.3.11
- earthengine-api 1.4.3 -> 1.5.9
- google-cloud-bigquery 3.29.0 -> 3.31.0
- bigquery-magics 0.8.1 -> 0.9.0
- NumPy 1.26.4 -> 2.0.2

Python package inclusions

- cuml-cu12 24.12.0

2025-01-13

- Released version 1.2.0 of the (Open in Colab Chrome Extension).
- Released minimizable comments with indicators next to cell.
- TPU v5e-1 Runtimes are now available for selection (tweet).
- GPU prices were decreased (tweet).

Python package upgrades

- accelerate 1.1.1 -> 1.2.1
- aiohttp 3.10.10 -> 3.11.11
- altair 4.2.2 -> 5.5.0
- bigframes 1.25.0 -> 1.29.0
- cmake 3.30.5 -> 3.31.2
- cvxpy 1.5.3 -> 3.6.0
- earthengine-api 1.2.0 -> 1.4.3
- folium 0.18.0 -> 0.19.3
- holidays 0.60 -> 0.63
- huggingface-hub 0.26.2 -> 0.27.0
- jsonpickle 3.4.2 -> 4.0.1
- kagglehub 0.3.3 -> 0.3.6
- keras 3.4.1 -> 3.5.0
- matplotlib 3.8.0 -> 3.10.0
- openai 1.54.3 -> 1.57.4
- pymc 5.18.0 -> 5.19.1
- safetensors 0.4.5 -> 0.5.0
- scikit-image 0.24.0 -> 0.25.0
- scikit-learn 1.5.2 -> 1.6.0
- sentence-transformers 3.2.1 -> 3.3.1
- tensorflow 2.17.0 -> 2.17.1
- torch 2.5.0 -> 2.5.1
- torchaudio 2.5.0 -> 2.5.1

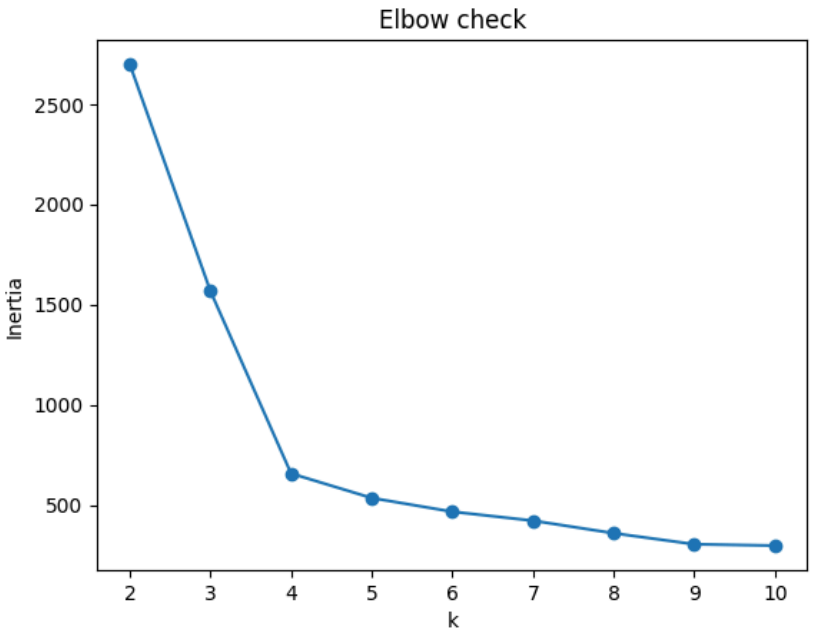

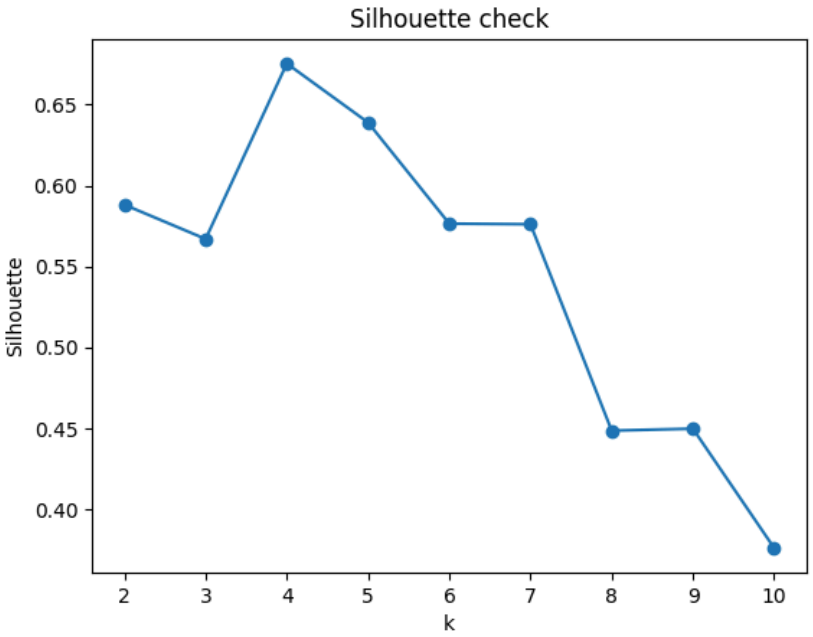

Regular K-Means scans the full dataset on every iteration, recalculating distances for every sample, so it usually yields tighter clusters with lower inertia and higher average silhouette values, but it is slower and needs the data to fit comfortably in memory. Mini-Batch K-Means speeds this up by updating the centroids with small random subsets drawn each step; this cuts computation time and memory needs dramatically while converging to a solution that is often close in quality, though its inertia is typically a bit higher because the noisier updates stop it exactly matching the full-data optimum.

Clustering is an unsupervised task: the algorithm receives no labels and must discover natural groupings based solely on similarity, producing clusters that may or may not align with any real-world categories. Classification, by contrast, is supervised; it learns a mapping from input features to known class labels

- Users can now import Gemini API keys from AI Studio into their user secrets, all in Colab (tweet).
- Increased limit to 1000 characters for requests to Gemini in Chat and Generate windows.
- Improved saving notebook to GitHub flow.
- Updated Gemini spark icon to be colorful.
- uv is pre-installed on the PATH for faster package installs.
- Fixed bugs
  - Dropdown text for GitHub repository not visible #4901.
  - Pre-installed California housing dataset README not correct #4862.
  - Backend execution error for scheduled notebook #4850.
  - Drive File Stream issues #3441.
  - Linking to the signup page does not preserve the authuser parameter.
  - Error messages in Gemini chat are not polished.
  - Clicking in Gemini chat feedback causes jitters the UI.
  - Hovering over a table of contents entry would show the menu icons for all entries.
  - Surveys display over open dialogs.
  - Playground mode banner not shown on mobile.

Python package upgrades

- accelerate 0.34.2 -> 1.1.1
- arviz 0.19.0 -> 0.20.0
- bigframes 1.18.0 -> 1.25.0
- bigquery-magics 0.2.0 -> 0.4.0
- bokeh 3.4.3 -> 3.6.1
- blosc 2.0.0 -> 2.7.1
- cloudpickle 2.2.1 -> 3.1.0
- cudf-cu12 24.4.1 -> 24.10.1
- dask 2024.8.0 -> 24.10.0
- debugpy 1.6.6 -> 1.8.0
- earthengine-api 1.0.0 -> 1.2.0
- folium 0.17.0 -> 0.18.0
- gscfs 2024.6.1 -> 2024.10.0
- geemap 0.34.3 -> 0.35.1
- holidays 0.57 -> 0.60
- huggingface-hub 0.24.7 -> 0.26.2
- kagglehub 0.3.0 -> 0.3.3
- lightgbm 4.4.0 -> 4.5.0
- lxml 4.9.4 -> 5.3.0
- matplotlib 3.7.1 -> 3.8.0
- mizani 0.11.4 -> 0.13.0
- networkx 3.3 -> 3.4.2
- nltk 3.8.1 -> 3.9.1
- pandas 2.1.4 -> 2.2.2
- pillow 10.4.0 -> 11.0.0
- plotnine 0.13.6 -> 0.14.1
- polars 1.6.0 -> 1.9.0
- protobuf 3.20.3 -> 4.25.5
- pyarrow 14.0.2 -> 17.0.0
- pydrive2 1.20.0 -> 1.21.1
- pymc 5.16.2 -> 5.18.0
- torch 2.4.1 -> 2.5.0
- torchaudio 2.4.1 -> 2.5.0

during training and later predicts those labels for unseen examples, so success is measured against the ground-truth classes.

From working with K-Means here I also learned that initializing centroids with the K-Means++ strategy, which purposely spreads them apart at the start, greatly lowers the chance of falling into a bad local minimum and often lets you reduce the number of random restarts. In addition, the silhouette score provides a quick, label-free way to judge cluster quality and pick a sensible value for k by balancing cohesion within clusters and separation between clusters.

• Clustering vs Classification

Clustering groups unlabeled data by similarity. There are no ground truth classes; the goal is to discover structure (unsupervised).

Classification assigns predefined class labels to inputs. It trains on labeled examples and predicts labels for new ones (supervised).

• one extra thing I learned from the notebook

K-Means++ initialization puts initial centroids far apart, reducing the risk of bad local minima and often letting you lower n_init.

- pygit2 1.16.0
- pyspark 3.5.3
- sentence-transformers 3.2.1
- timm 1.0.11
- wandb 0.18.6

Library and driver upgrades

- drivefs upgraded from 89.0.2 to 98.0.0

2024-09-23

- Improved code snippet search
- Updated Marketplace image and public local runtime container
- Improved the look-and-feel of interactive form dropdowns and checkboxes
- Fixed bugs
  - activating the skip link caused the notebook to scroll out of view
  - toggling a checkbox too much caused the page to crash
  - lightning fast drags could cause orphaned tabs
  - custom widgets snippet would show for local runtimes

Python package upgrades

- accelerate 0.32.1 -> 0.34.2
- arviz 0.18.0 -> 0.19
- autograd 1.6.2 -> 1.7.0
- bigframes 1.14.0 -> 1.18.0
- dask 2024.7.1 -> 2024.8.0
- distributed 2024.7.1 -> 2024.8.0
- duckdb 0.10.3 -> 1.1.0
- earthengine-api 0.1.416 -> 1.0.0
- flax 0.8.4 -> 0.8.5
- gdown 5.1.0 -> 5.2.0
- geemap 0.33.1 -> 0.34.3
- geopandas 0.14.4 -> 1.0.1
- google-cloud-aiplatform 1.59.0 -> 1.67.1
- google-cloud-bigquery-storage 2.25.0 -> 2.26.0
- holidays 0.54 -> 0.57
- huggingface-hub 0.23.5 -> 0.24.7
- ibis-framework 8.0.0 -> 9.2.0
- jax 0.4.26 -> 0.4.33
- jaxlib 0.4.26 -> 0.4.33
- kagglehub 0.2.9 -> 0.3.0
- lightgbm 4.4.0 -> 4.5.0
- matplotlib-venn 0.11.10 -> 1.1.1
- mizani 0.9.3 -> 0.11.4
- Pillow 9.4.0 -> 10.4.0
- plotly 5.15.0 -> 5.24.1
- plotnine 0.12.4 -> 0.13.6
- polars 0.20.2 -> 1.6.0
- progressbar2 4.2.0 -> 4.5.0
- PyDrive2 1.6.3 -> 1.20.0
- pymc 5.10.4 -> 5.16.2
- pytensor 2.18.6 -> 2.25.4
- scikit-image 0.23.2 -> 0.24.0
- scikit-learn 1.3.2 -> 1.5.2
- torch 2.3.1 -> 2.4.1
- torchaudio 2.3.1 -> 2.4.1
- torchvision 0.18.1 -> 0.19.1
- transformers 4.42.4 -> 4.44.2
- urllib3 2.0.7 -> 2.2.3
- xarray 2024.6.0 -> 2024.9.0

relevant sources

- Added a new "Create Gemini API key" link in the user secrets panel
- Added a new "Gemini: Creating a prompt" snippet and touched up the existing "Gemini: Connecting to Gemini" snippet
- Added the ability to specify custom placeholder text for various interactive form params (see examples)
- Keyboard navigation a11y improvements to comments UI
- Various minor rendering improvements to interactive forms UI
- A11y improvements for the run button and header
- Updated tooltip styling
- A11y improvements for the file browser's disk usage bar
- On mobile, tooltips now trigger on long press
- On mobile, release notes updates will no longer display automatically
- Python package upgrades
  - astropy 5.3.4 -> 6.1.2
  - bigframes 1.11.1 -> 1.14.0
  - bokeh 3.3.4 -> 3.4.3
  - dask 2023.8.1 -> 2024.7.1
  - earthengine-api 0.1.412 -> 0.1.416
  - geopandas 0.13.2 -> 0.14.4
  - kagglehub 0.2.8 -> 0.2.9
  - keras 2.15.0 -> 3.4.1
  - lightgbm 4.1.0 -> 4.4.0
  - malloy 2023.1067 -> 2024.1067
  - numba 0.58.1 -> 0.60.0
  - numpy 1.25.2 -> 1.26.4
  - opencv-python 4.8.0.76 -> 4.10.0.84
  - pandas 2.0.3 -> 2.1.4
  - pandas-gbq 0.19.2 -> 0.23.1
  - panel 1.3.8 -> 1.4.5
  - requests 2.31.0 -> 2.32.3
  - scikit-learn 1.2.2. -> 1.3.2
  - scipy 1.11.4 -> 1.13.1
  - tensorboard 2.15.2 -> 2.17.0
  - tensorflow 2.15.0 -> 2.17.0
  - tf-keras 2.15.1 -> 2.17.0
  - xarray 2023.7.0 -> 2024.6.0
  - xgboost 2.0.3 -> 2.1.1
- Python package inclusions
  - einops 0.8.0

2024-07-22

- You can now embed Google sheets directly into Colab to streamline interactions with data with InteractiveSheet.

  Example:

  ```
  from google.colab import sheets
  sh = sheets.InteractiveSheet()
  df = sh.as_df()
  ```

- Fixed multiple rendering bugs in cell editors with wide text content (i.e. text is no longer

- Fixed multiple scrollbar bugs in the user secrets panel
- Added the ability for workspace admin to purchase Colab Pro and Pro+ Subscriptions for users
- Fixed bug where user secrets couldn't be moved to a tab
- Fixed several focus management accessibility issues in tabs, the table of contents, the left toolbar, and the run button
- Fixed bug where overflowing cells may be omitted when pasting from Google Sheets
- Fixed bug where the generate code button did not activate on touch
- Python package upgrades
  - bigframes 1.9.0 -> 1.11.1
  - cvxpy 1.3.4 -> 1.5.2
  - earthengine-api 0.1.408 -> 0.1.412
  - google-api-core 2.11.1 -> 2.19.1
  - google-api-python-client 2.84.0 -> 2.137.0
  - google-cloud-aiplatform 1.56.0 -> 1.59.0
  - google-cloud-bigquery 3.21.0 -> 3.25.0
  - google-cloud-core 2.3.3 -> 2.4.1
  - google-cloud-datastore 2.15.2 -> 2.19.0
  - google-cloud-firestore 2.11.1 -> 2.16.1
  - google-cloud-functions 1.13.3 -> 1.16.4
  - google-generativeai 0.5.4 -> 0.7.2
  - kagglehub 0.2.5 -> 0.2.8
  - pip 23.1.2 -> 24.1.2
  - setuptools 67.7.2 -> 71.0.4
  - sympy 1.12.1 -> 1.13.1
  - torch 2.3.0 -> 2.3.1
  - transformers 4.41.2 -> 4.42.4
- Python package inclusions
  - accelerate 0.32.1

2024-06-18

- Inline AI completions are now available to users on the free-of-charge tier
- Reduced latency for LSP and terminal connections
- Improved quality of inline completions
- Visual improvements to switch controls across Colab
- Various bug fixes, performance and a11y improvements to the user secrets panel
- Improved tooltip UX behavior
- Improved behavior when copying data from Google Sheets and pasting in Colab
- Scroll to cell fixes for single tabbed view and jump to cell command
- Improved tab header behavior
- A11y improvements for notebook-focused cells
- Python package upgrades
  - torch 2.2.1 -> 2.3.0
  - torchaudio 2.2.1 -> 2.3.0
  - torchvision 0.17.1 -> 0.18.0
  - torchtext 0.17.1 -> 0.18.0

shortcut "Ctrl/⌘ + ."

- Python package upgrades
  - bigframes 1.0.0 -> 1.5.0
  - google-cloud-aiplatform 1.47.0 -> 1.51.0
  - jax[tpu] 0.4.23 -> 0.4.26
- Python package inclusions
  - cudf 24.4.1

2024-04-15

- TPU v2 runtime is now available
- L4 runtime is now available for paid users
- New distributed fine-tuning Gemma tutorial on TPUs (GitHub)
- Symbol rename is now supported with keyboard shortcut F2
- Fixed bug causing inability to re-upload deleted files
- Fixed breaking bug in colabtools %upload_files_async
- Added syntax highlighting to %%writefile cells
- Cuda dependencies that come with Torch are cached for faster downloads for packages that require Torch and its dependencies (GitHub issue)
- Python package upgrades
  - bigframes 0.24.0 -> 1.0.0
  - duckdb 0.9.2 -> 0.10.1
  - google-cloud-aiplatform 1.43.0 -> 1.47.0
  - jax 0.4.23 -> 0.4.26

2024-03-13

- Fixed bug that sometimes caused UserSecrets to move / disappear
- Improved messaging for mounting drive in an unsupported environment (GitHub issue)
- Python package upgrades
  - torch 2.1.0 -> 2.2.1
  - torchaudio 2.1.0 -> 2.2.1
  - torchvision 0.16.0 -> 0.17.1
  - torchtext 0.16.0 -> 0.17.1
  - PyMC 5.7.2 -> 5.10.4
  - BigFrames 0.21.0 -> 0.24.0
  - google-cloud-aiplatform 1.42.1 -> 1.43.0
  - tornado 6.3.2 -> 6.3.3

2024-02-21

- Try out Gemma on Colab!
- Allow unicode in form text inputs
- Display documentation and link to source when displaying functions
- Display image-like ndarrays as images
- Improved UX around quick charts and execution error suggestions
- Released Marketplace image for the month of February (GitHub issue)
- Python package upgrades

- pyarrow 10.0.1 -> 14.0.2

2024-01-29

- New Kaggle Notebooks <> Colab updates! Now you can:
  - Import directly from Colab without having to download/re-upload
  - Upload via link, by pasting Google Drive or Colab URLs
  - Export & run Kaggle Notebooks on Colab with 1 click
- Try these notebooks that talk to Gemini:
  - Gemini and Stable Diffusion
  - Learning with Gemini and ChatGPT
  - Talk to Gemini with Google's Speech to Text API
  - Sell lemonade with Gemini and Sheets
  - Generate images with Gemini and Vertex
- Python package upgrades
  - google-cloud-aiplatform 1.38.1 -> 1.39.0
  - bigframes 0.18.0 -> 0.19.2
  - polars 0.17.3 -> 0.20.2
  - gdown 4.6.6 -> 4.7.3 (GitHub issue)
  - tensorflow-hub 0.15.0 -> 0.16.0
  - flax 0.7.5 -> 0.8.0
- Python package inclusions
  - sentencepiece 0.1.99

2024-01-08

- Avoid nested scrollbars for large outputs by using `google.colab.output.no_vertical_scr` Example notebook
- Fix bug where downloading models from Hugging Face could freeze
- Python package upgrades
  - huggingface-hub 0.19.4 -> 0.20.2
  - bigframes 0.17.0 -> 0.18.0

2023-12-18

- Expanded access to AI coding has arrived in Colab across 175 locales for all tiers of Colab users
- Improvements to display of ML-based inline completions (for eligible Pro/Pro+ users)
- Started a series of notebooks highlighting Gemini API capabilities
- Enable ⌘/Ctrl+L to select the full line in an editor
- Fixed bug where we weren't correctly formatting output from multiple execution results
- Python package upgrades
  - CUDA 11.8 to CUDA 12.2
  - tensorflow 2.14.0 -> 2.15.0
  - tensorboard 2.14.0 -> 2.15.0
  - keras 2.14.0 -> 2.15.0
  - Nvidia drivers 525.105.17 -> 535.104.05
  - tensorflow-gcs-config 2.14.0 -> 2.15.0
  - bigframes 0.13.0 -> 0.17.0

2023-11-27

- Removed warning when calling `await` to make it render as code
- Added "Run selection" to the cell context menu
- Added highlighting for the `%%python` cell magic
- Launched AI coding features for Pro/Pro+ users in more locales
- Python package upgrades
  - bigframes 0.12.0 -> 0.13.0
- Python package inclusions
  - transformers 4.35.2
  - google-generativeai 0.2.2

2023-11-08

- Launched Secrets, for safe storage of private keys on Colab (tweet)
- Fixed issue where TensorBoard would not load (#3990)
- Python package upgrades
  - lightgbm 4.0.0 -> 4.1.0
  - bigframes 0.10.0 -> 0.12.0
  - bokeh 3.2.2 -> 3.3.0
  - duckdb 0.8.1 -> 0.9.1
  - numba 0.56.4 -> 0.58.1
  - tweepy 4.13.0 -> 4.14.0
  - jax 0.4.16 -> 0.4.20
  - jaxlib 0.4.16 -> 0.4.20

2023-10-23

- Updated the **Open notebook** dialog for better usability and support for smaller screen sizes
- Added smart paste support for data from Google Sheets for R notebooks
- Enabled showing release notes in a tab
- Launched AI coding features for Pro/Pro+ users in Australia Canada India and Japan (tweet)
- Python package upgrades
  - earthengine-api 0.1.357 -> 0.1.375
  - flax 0.7.2 -> 0.7.4
  - geemap 0.27.4 -> 0.28.2
  - jax 0.4.14 -> 0.4.16
  - jaxlib 0.4.14 -> 0.4.16
  - keras 2.13.1 -> 2.14.0
  - tensorboard 2.13.0 -> 2.14.1
  - tensorflow 2.13.0 -> 2.14.0
  - tensorflow-gcs-config 2.13.0 -> 2.14.0
  - tensorflow-hub 0.14.0 -> 0.15.0
  - tensorflow-probability 0.20.1 -> 0.22.0
  - torch 2.0.1 -> 2.1.0
  - torchaudio 2.0.2 -> 2.1.0
  - torchtext 0.15.2 -> 0.16.0
  - torchvision 0.15.2 -> 0.16.0
  - xgboost 1.7.6 -> 2.0.0
- Python package inclusions
  - bigframes 0.10.0
  - malloy 2023.1056

forms
- Improved rendering of the notebook when printing
- Python package upgrades
  - tensorflow 2.12.0 -> 2.13.0
  - tensorboard 2.12.3 -> 2.13.0
  - keras 2.12.0 -> 2.13.1
  - tensorflow-gcs-config 2.12.0 -> 2.13.
  - scipy 1.10.1-> 1.11.2
  - cython 0.29.6 -> 3.0.2
- Python package inclusions
  - geemap 0.26.0

2023-08-18

- Added "Change runtime type" to the menu in the connection button
- Improved auto-reconnection to an already running notebook (#3764)
- Increased the specs of our highmem machines for Pro users
- Fixed `add-apt-repository` command on Ubuntu 22.04 runtime (#3867)
- Python package upgrades
  - bokeh 2.4.3 -> 3.2.2
  - cmake 3.25.2 -> 3.27.2
  - cryptography 3.4.8 -> 41.0.3
  - dask 2022.12.1 -> 2023.8.0
  - distributed 2022.12.1 -> 2023.8.0
  - earthengine-api 0.1.358 -> 0.1.364
  - flax 0.7.0 -> 0.7.2
  - ipython-sql 0.4.0 -> 0.5.0
  - jax 0.4.13 -> 0.4.14
  - jaxlib 0.4.13 -> 0.4.14
  - lightgbm 3.3.5 -> 4.0.0
  - mkl 2019.0 -> 2023.2.0
  - notebook 6.4.8 -> 6.5.5
  - numpy 1.22.4 -> 1.23.5
  - opencv-python 4.7.0.72 -> 4.8.0.76
  - pillow 8.4.0 -> 9.4.0
  - plotly 5.13.1 -> 5.15.0
  - prettytable 0.7.2 -> 3.8.0
  - pytensor 2.10.1 -> 2.14.2
  - spacy 3.5.4 -> 3.6.1
  - statsmodels 0.13.5 -> 0.14.0
  - xarray 2022.12.0 -> 2023.7.0
- Python package inclusions
  - PyDrive2 1.6.3

2023-07-21

- Launched auto-plotting for dataframes, available using the chart button that shows up alongside datatables (post)

- Added a menu to the table of contents to support running a section or collapsing/expanding sections (post)

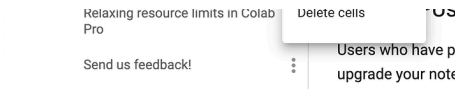

- Added an option to automatically run the first cell or section, available under Edit -> Notebook settings (post)

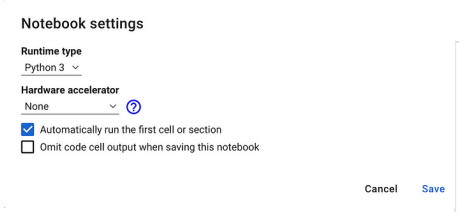

- Launched Pro/Pro+ to Algeria, Argentina, Chile, Ecuador, Egypt, Ghana, Kenya, Malaysia, Nepal, Nigeria, Peru, Rwanda, Saudi Arabia, South Africa, Sri Lanka, Tunisia, and Ukraine (tweet)
- Added a command, "Toggle tab moves focus" for toggling tab trapping in the editor (Tools -> Command palette, "Toggle tab moves focus")
- Fixed issue where `files.upload()` was sometimes returning an incorrect filename (#1550)
- Fixed f-string syntax highlighting bug (#3802)
- Disabled ambiguous characters highlighting for commonly used LaTeX characters (#3648)
- Upgraded Ubuntu from 20.04 LTS to 22.04 LTS
- Updated the Colab Marketplace VM image
- Python package upgrades:
  - autograd 1.6.1 -> 1.6.2
  - drivefs 76.0 -> 77.0
  - flax 0.6.11 -> 0.7.0
  - earthengine-api 0.1.357 -> 0.1.358
  - GDAL 3.3.2->3.4.3
  - google-cloud-bigquery-storage 2.20.0 -> 2.22.2
  - gspread-dataframe 3.0.8 -> 3.3.1
  - holidays 0.27.1 -> 0.29
  - jax 0.4.10 -> jax 0.4.13
  - jaxlib 0.4.10 -> jax 0.4.13
  - jupyterlab-widgets 3.0.7 -> 3.0.8
  - nbformat 5.9.0 -> 5.9.1
  - opencv-python-headless 4.7.0.72 -> 4.8.0.74
  - pygame 2.4.0 -> 2.5.0
  - spacy 3.5.3 -> 3.5.4
  - SQLAlchemy 2.0.16 -> 2.0.19
  - tabulate 0.8.10 -> 0.9.0
  - tensorflow-hub 0.13.0 -> 0.14.0

2023-06-23

- Launched AI coding features to subscribed users starting with Pro+ users in the US (tweet, post)
- Added the Kernel Selector in the Notebook Settings (tweet)
- Fixed double space trimming issue in markdown #3766

- flax 0.6.9 -> 0.6.11
- google-cloud-bigquery 3.9.0 -> 3.10.0
- google-cloud-bigquery-storage 2.19.1 -> 2.20.0
- grpcio 1.54.0 -> 1.56.0
- holidays 0.25 -> 0.27.1
- nbformat 5.8.0 -> 5.9.0
- prophet 1.1.3 -> 1.1.4
- pydata-google-auth 1.7.0 -> 1.8.0
- spacy 3.5.2 -> 3.5.3
- tensorboard 2.12.2 -> 2.12.3
- xgboost 1.7.5 -> 1.7.6

- Python package inclusions:
  - gcsfs 2023.6.0
  - geopandas 0.13.2
  - google-cloud-bigquery-connection 1.12.0
  - google-cloud-functions 1.13.0
  - grpc-google-iam-v1 0.12.6
  - multidict 6.0.4
  - tensorboard-data-server 0.7.1

2023-06-02

- Released the new site colab.google
- Published Colab's Docker runtime image to us-docker.pkg.dev/colab-images/public/runtime (tweet, instructions)
- Launched support for Google children accounts (tweet)
- Launched DagsHub integration (tweet, post)
- Upgraded to Monaco Editor Version 0.37.1
- Fixed various Vim keybinding bugs
- Fixed issue where the N and P letters sometimes couldn't be typed (#3664)
- Fixed rendering support for compositional inputs (#3660, #3679)
- Fixed lag in notebooks with lots of cells (#3676)
- Improved support for R by adding a Runtime type notebook setting (Edit -> Notebook settings)
- Improved documentation for connecting to a local runtime (Connect -> Connect to a local runtime)
- Python package updates:
  - holidays 0.23 -> 0.25
  - jax 0.4.8 -> 0.4.10
  - jaxlib 0.4.8 -> 0.4.10
  - pip 23.0.1 -> 23.1.2
  - tensorflow-probability 0.19.0 -> 0.20.1
  - torch 2.0.0 -> 2.0.1
  - torchaudio 2.0.1 -> 2.0.2
  - torchdata 0.6.0 -> 0.6.1
  - torchtext 0.15.1 -> 0.15.2
  - torchvision 0.15.1 -> 0.15.2
  - tornado 6.2 -> 6.3.1

2023-05-05

- Released GPU type selection for paid users, allowing them to choose a preferred NVidia GPU
- Upgraded R from 4.2.3 to 4.3.0
- Upgraded Python from 3.9.16 to 3.10.11
- Python package updates:
  - attrs 22.2.0 -> attrs 23.1.0

2023-04-14

- Python package updates:
  - google-api-python-client 2.70.0 -> 2.84.0
  - google-auth-oauthlib 0.4.6 -> 1.0.0
  - google-cloud-bigquery 3.4.2 -> 3.9.0
  - google-cloud-datastore 2.11.1 -> 2.15.1
  - google-cloud-firestore 2.7.3 -> 2.11.0
  - google-cloud-language 2.6.1 -> 2.9.1
  - google-cloud-storage 2.7.0 -> 2.8.0
  - google-cloud-translate 3.8.4 -> 3.11.1
  - networkx 3.0 -> 3.1
  - notebook 6.3.0 -> 6.4.8
  - jax 0.4.7 -> 0.4.8
  - pandas 1.4.4 -> 1.5.3
  - spacy 3.5.1 -> 3.5.2
  - SQLAlchemy 1.4.47 -> 2.0.9
  - xgboost 1.7.4 -> 1.7.5

2023-03-31

- Improve bash ! syntax highlighting (GitHub issue)
- Fix bug where VIM keybindings weren't working in the file editor
- Upgraded R from 4.2.2 to 4.2.3
- Python package updates:
  - arviz 0.12.1 --> 0.15.1
  - astropy 4.3.1 --> 5.2.2
  - dopamine-rl 1.0.5 --> 4.0.6
  - gensim 3.6.0 --> 4.3.1
  - ipykernel 5.3.4 -> 5.5.6
  - ipython 7.9.0 -> 7.34.0
  - jax 0.4.4 -> 0.4.7
  - jaxlib 0.4.4 -> 0.4.7
  - jupyter_core 5.2.0 -> 5.3.0
  - keras 2.11.0 -> 2.12.0
  - lightgbm 2.2.3 -> 3.3.5
  - matplotlib 3.5.3 -> 3.7.1
  - nltk 3.7 -> 3.8.1
  - opencv-python 4.6.0.66 -> 4.7.0.72
  - plotly 5.5.0 -> 5.13.1
  - pymc 4.1.4 -> 5.1.2
  - seaborn 0.11.2 -> 0.12.2
  - spacy 3.4.4 -> 3.5.1
  - sympy 1.7.1 -> 1.11.1
  - tensorboard 2.11.2 -> 2.12.0
  - tensorflow 2.11.0 -> 2.12.0
  - tensorflow-estimator 2.11.0 -> 2.12.0
  - tensorflow-hub 0.12.0 -> 0.13.0
  - torch 1.13.1 -> 2.0.0
  - torchaudio 0.13.1 -> 2.0.1
  - torchtext 0.14.1 -> 0.15.1
  - torchvision 0.14.1 -> 0.15.1

2023-03-10

- Added the Colab editor shortcuts example notebook
- Fixed triggering of @-mention and email autocomplete for large comments (GitHub issue)
- Added View Resources to the Runtime menu
- Made file viewer images fit the view by default, resizing to original size on click

- beautifulsoup4 4.6.3 -> 4.9.3
- bokeh 2.3.3 -> 2.4.3
- debugpy 1.0.0 -> 1.6.6
- Flask 1.1.4 -> 2.2.3
- jax 0.3.25 -> 0.4.4
- jaxlib 0.3.25 -> 0.4.4
- Jinja2 2.11.3 -> 3.1.2
- matplotlib 3.2.2 -> 3.5.3
- nbconvert 5.6.1 -> 6.5.4
- pandas 1.3.5 -> 1.4.4
- pandas-datareader 0.9.0 -> 0.10.0
- pandas-profiling 1.4.1 -> 3.2.0
- Pillow 7.1.2 -> 8.4.0
- plotnine 0.8.0 -> 0.10.1
- scikit-image 0.18.3 -> 0.19.3
- scikit-learn 1.0.2 -> 1.2.2
- scipy 1.7.3 -> 1.10.1
- setuptools 57.4.0 -> 63.4.3
- sklearn-pandas 1.8.0 -> 2.2.0
- statsmodels 0.12.2 -> 0.13.5
- urllib3 1.24.3 -> 1.26.14
- Werkzeug 1.0.1 -> 2.2.3
- wrapt 1.14.1 -> 1.15.0
- xgboost 0.90 -> 1.7.4
- xlrd 1.2.0 -> 2.0.1

2023-02-17

- Show graphs of RAM and disk usage in notebook toolbar
- Copy cell links directly to the clipboard instead of showing a dialog when clicking on the link icon in the cell toolbar
- Updated the Colab Marketplace VM image
- Upgraded CUDA to 11.6.2 and cuDNN to 8.4.0.27
- Python package updates:
  - tensorflow 2.9.2 -> 2.11.0
  - tensorboard 2.9.1 -> 2.11.2
  - keras 2.9.0 -> 2.11.0
  - tensorflow-estimator 2.9.0 -> 2.11.0
  - tensorflow-probability 0.17.0 -> 0.19.0
  - tensorflow-gcs-config 2.9.0 -> 2.11.0
  - earthengine-api 0.1.339 -> 0.1.341
  - flatbuffers 1.12 -> 23.1.21
  - platformdirs 2.6.2 -> 3.0.0
  - pydata-google-auth 1.6.0 -> 1.7.0
  - python-utils 3.4.5 -> 3.5.2
  - tenacity 8.1.0 -> 8.2.1
  - tifffile 2023.1.23.1 -> 2023.2.3
  - notebook 5.7.16 -> 6.3.0
  - tornado 6.0.4 -> 6.2
  - aiohttp 3.8.3 -> 3.8.4
  - charset-normalizer 2.1.1 -> 3.0.1
  - fastai 2.7.0 -> 2.7.1
  - soundfile 0.11.0 -> 0.12.1
  - typing-extensions 4.4.0 -> 4.5.0
  - widgetsnbextension 3.6.1 -> 3.6.2
  - pydantic 1.10.4 -> 1.10.5
  - zipp 3.12.0 -> 3.13.0
  - numpy 1.21.6 -> 1.22.4
  - drivefs 66.0 -> 69.0
  - gdal 3.0.4 -> 3.3.2 GitHub issue
- Added libudunits2-dev for smoother R package installs GitHub issue

- absl-py 1.3.0 -> 1.4.0
- bleach 5.0.1 -> 6.0.0
- cachetools 5.2.1 -> 5.3.0
- cmdstanpy 1.0.8 -> 1.1.0
- dnspython 2.2.1 -> 2.3.0
- fsspec 2022.11.0 -> 2023.1.0
- google-cloud-bigquery-storage 2.17.0 -> 2.18.1
- holidays 0.18 -> 0.19
- jupyter-core 5.1.3 -> 5.2.0
- packaging 21.3 -> 23.0
- prometheus-client 0.15.0 -> 0.16.0
- pyct 0.4.8 -> 0.5.0
- pydata-google-auth 1.5.0 -> 1.6.0
- python-slugify 7.0.0 -> 8.0.0
- sqlalchemy 1.4.46 -> 2.0.0
- tensorflow-io-gcs-filesystem 0.29.0 -> 0.30.0
- tifffile 2022.10.10 -> 2023.1.23.1
- zipp 3.11.0 -> 3.12.0
- Pinned sqlalchemy to version 1.4.46

2023-01-12

- Added support for @-mention and email autocomplete in comments
- Improved errors when GitHub notebooks can't be loaded
- Increased color contrast for colors used for syntax highlighting in the code editor
- Added terminal access for custom GCE VM runtimes
- Upgraded Ubuntu from 18.04 LTS to 20.04 LTS (GitHub issue)
- Python package updates:
  - GDAL 2.2.2 -> 2.2.3.
  - NumPy from 1.21.5 to 1.21.6.
  - attrs 22.1.0 -> 22.2.0
  - chardet 3.0.4 -> 4.0.0
  - cloudpickle 1.6.0 -> 2.2.0
  - filelock 3.8.2 -> 3.9.0
  - google-api-core 2.8.2 -> 2.11.0
  - google-api-python-client 1.12.11 -> 2.70.0
  - google-auth-httplib2 0.0.3 -> 0.1.0
  - google-cloud-bigquery 3.3.5 -> 3.4.1
  - google-cloud-datastore 2.9.0 -> 2.11.0
  - google-cloud-firestore 2.7.2 -> 2.7.3
  - google-cloud-storage 2.5.0 -> 2.7.0
  - holidays 0.17.2 -> holidays 0.18
  - importlib-metadata 5.2.0 -> 6.0.0
  - networkx 2.8.8 -> 3.0
  - opencv-python-headless 4.6.0.66 -> 4.7.0.68
  - pip 21.1.3 -> 22.04
  - pip-tools 6.2.0 -> 6.6.2
  - prettytable 3.5.0 -> 3.6.0
  - requests 2.23.0 -> 2.25.1
  - termcolor 2.1.1 -> 2.2.0
  - torch 1.13.0 -> 1.13.1
  - torchaudio 0.13.0 -> 0.13.1
  - torchtext 0.14.0-> 0.14.1
  - torchvision 0.14.0 -> 0.14.1

2022-12-06

  - torchaudio from 0.12.1 to 0.13.0
  - torchvision from 0.13.1 to 0.14.0
  - torchtext from 0.13.1 to 0.14.0
  - xlrd from 1.1.0 to 1.2.0
  - DriveFS from 62.0.1 to 66.0.3
- Made styling of markdown tables in outputs match markdown tables in text cells
- Improved formatting for empty interactive table rows
- Fixed syntax highlighting for variables with names that contain Python keywords (GitHub issue)

2022-11-11

- Added more dark editor themes for Monaco (when in dark mode, "Editor colorization" appears as an option in the Editor tab of the Tools → Settings dialog)
- Fixed bug where collapsed forms were deleted on mobile GitHub issue
- Python package updates:
  - rpy2 from 3.4.0 to 3.5.5 (GitHub issue)
  - notebook from 5.5.0 to 5.7.16
  - tornado from 5.1.1 to 6.0.4
  - tensorflow_probability from 0.16.0 to 0.17.0
  - pandas-gbq from 0.13.3 to 0.17.9
  - protobuf from 3.17.3 to 3.19.6
  - google-api-core[grpc] from 1.31.5 to 2.8.2
  - google-cloud-bigquery from 1.21.0 to 3.3.5
  - google-cloud-core from 1.0.1 to 2.3.2
  - google-cloud-datastore from 1.8.0 to 2.9.0
  - google-cloud-firestore from 1.7.0 to 2.7.2
  - google-cloud-language from 1.2.0 to 2.6.1
  - google-cloud-storage from 1.18.0 to 2.5.0
  - google-cloud-translate from 1.5.0 to 3.8.4

2022-10-21

- Launched a single-click way to get from BigQuery to Colab to further explore query results (announcement)
- Launched Pro, Pro+, and Pay As You Go to 19 additional countries: Austria, Belgium, Bulgaria, Croatia, Cyprus, Czechia, Denmark, Estonia, Finland, Greece, Hungary, Latvia, Lithuania, Norway, Portugal, Romania, Slovakia, Slovenia, and Sweden (tweet)
- Updated jax from 0.3.17 to 0.3.23, jaxlib from 0.3.15 to 0.3.22, TensorFlow from 2.8.2 to 2.9.2, CUDA from 11.1 to 11.2, and cuDNN from 8.0 to 8.1 (backend-info)
- Added a `readonly` option to `drive.mount`
- Fixed bug where Xarray was not working (GitHub issue)
- Modified Markdown parsing to ignore block quote symbol within MathJax (GitHub issue)

2022-09-30

- Launched Pay As You Go, allowing premium GPU access without requiring a subscription

- Created a GitHub repo, backend-info, with the latest apt-list.txt and pip-freeze.txt files for the Colab runtime (GitHub issue)
- Added `files.upload_file(filename)` to upload a file from the browser to the runtime with a specified filename

2022-09-16

- Upgraded pymc from 3.11.0 to 4.1.4, jax from 0.3.14 to 0.3.17, jaxlib from 0.3.14 to 0.3.15, fsspec from 2022.8.1 to 2022.8.2
- Modified our save flow to avoid persisting Drive filenames as titles in notebook JSON
- Updated our Terms of Service
- Modified the `Jump to Cell` command to locate the cursor at the end of the command palette input (`Jump to cell` in Tools → Command palette in a notebook with section headings)
- Updated the styling of the Drive notebook comment UI
- Added support for terminating your runtime from code: `python from google.colab import runtime runtime.unassign()`
- Added regex filter support to the Recent notebooks dialog
- Inline google.colab.files.upload JS to fix `files.upload()` not working (GitHub issue)

2022-08-26

- Upgraded PyYAML from 3.13 to 6.0 (GitHub issue), drivefs from 61.0.3 to 62.0.1
- Upgraded TensorFlow from 2.8.2 to 2.9.1 and ipywidgets from 7.7.1 to 8.0.1 but rolled both back due to a number of user reports (GitHub issue, GitHub issue)
- Stop persisting inferred titles in notebook JSON (GitHub issue)
- Fix bug in background execution which affected some Pro+ users (GitHub issue)
- Fix bug where `Download as .py` incorrectly handled text cells ending in a double quote
- Fix bug for Pro and Pro+ users where we weren't honoring the preference (Tools → Settings) to use a temporary scratch notebook as the default landing page
- Provide undo/redo for scratch cells
- When writing ipynb files, serialize empty multiline strings as `[]` for better consistency with JupyterLab

2022-08-11

- Upgraded ipython from 5.5.0 to 7.9.0, fbprophet 0.7 to prophet 1.1, tensorflow-datasets from 4.0.1 to 4.6.0, drivefs from 60.0.2 to 61.0.3, pytorch from 1.12.0 to 1.12.1, numba from 0.51 to 0.56, and lxml from 4.2.0 to 4.9.1
- Loosened our `requests` version requirement (GitHub issue)
- Removed support for TensorFlow 1
- Added Help → Report Drive abuse for Drive notebooks
- Fixed indentation for Python lines ending in `[`

- Modified forms to use a value of None instead of causing a parse error when clearing raw and numeric-typed form fields

2022-07-22

- Update scipy from 1.4.1 to 1.7.3, drivefs from 59.0.3 to 60.0.2, pytorch from 1.11 to 1.12, jax & jaxlib from 0.3.8 to 0.3.14, opencv-python from 4.1.2.30 to 4.6.0.66, spaCy from 3.3.1 to 3.4.0, and dlib from 19.18.0 to 19.24.0
- Fix `Open in tab` doc link which was rendering incorrectly (GitHub issue)
- Add a preference for the default tab orientation to the Site section of the settings menu under Tools → Settings
- Show a warning for `USE_AUTH_EPHEM` usage when running authenticate_user on a TPU runtime (code)

2022-07-01

- Add a preference for code font to the settings menu under Tools → Settings
- Update drivefs from 58.0.3 to 59.0.3 and spacy from 2.2.4 to 3.3.1
- Allow display_data and execute_result text outputs to wrap, matching behavior of JupyterLab (does not affect stream outputs/print statements).
- Improve LSP handling of some magics, esp. %%writefile (GitHub issue).
- Add a FAQ entry about the mount Drive button behavior and include link buttons for each FAQ entry.
- Fix bug where the notebook was sometimes hidden behind other tabs on load when in single pane view.
- Fix issue with inconsistent scrolling when an editor is in multi-select mode.
- Fix bug where clicking on a link in a form would navigate away from the notebook
- Show a confirmation dialog before performing Replace all from the Find and replace pane.

2022-06-10

- Update drivefs from 57.0.5 to 58.0.3 and tensorflow from 2.8.0 to 2.8.2
- Support more than 100 repos in the GitHub repo selector shown in the open dialog and the clone to GitHub dialog
- Show full notebook names on hover in the open dialog
- Improve the color contrast for links, buttons, and the `ipywidgets.Accordion` widget in dark mode

2022-05-20

- Support URL params for linking to some common pref settings: force_theme=dark, force_corgi_mode=1, force_font_size=14. Params forced by URL are not persisted unless saved using Tools → Settings.
- Add a class `markdown-google-sans` to allow Markdown to render in Google Sans
- Update monaco-vim from 0.1.19 to 0.3.4

- Improved rendering of filter options in an interactive table
- Added git-lfs to the base image
- Updated torch from 1.10.0 to 1.11.0, jupyter-core from 4.9.2 to 4.10.0, and cmake from 3.12.0 to 3.22.3
- Added more details to our FAQ about unsupported uses (using proxies, downloading torrents, etc.)
- Fixed issue with apt-get dependencies

2022-04-15

- Add an option in the file browser to show hidden files.
- Upgrade gdown from 4.2.0 to 4.4.0, google-api-core[grpc] from 1.26.0 to 1.31.5, and pytz from 2018.4 to 2022.1

2022-03-25

- Launched Pro/Pro+ to 12 additional countries: Australia, Bangladesh, Colombia, Hong Kong, Indonesia, Mexico, New Zealand, Pakistan, Philippines, Singapore, Taiwan, and Vietnam
- Added `google.colab.auth.authenticate_serv` to support using Service Account keys
- Update jax from 0.3.1 to 0.3.4 & jaxlib from 0.3.0 to 0.3.2
- Fixed an issue with Twitter previews of notebooks shared as GitHub Gists

2022-03-10

- Launched Pro/Pro+ to 10 new countries: Ireland, Israel, Italy, Morocco, the Netherlands, Poland, Spain, Switzerland, Turkey, and the United Arab Emirates
- Launched support for scheduling notebooks for Pro+ users
- Fixed bug in interactive datatables where filtering by number did not work
- Finished removing the python2 kernelspec

2022-02-25

- Made various accessibility improvements to the header
- Fix bug with forms run:auto where a form field change would trigger multiple runs
- Minor updates to the bigquery example notebook and snippet
- Include background execution setting in the sessions dialog for Pro+ users
- Update tensorflow-probability from 0.15 to 0.16
- Update jax from 0.2.25 to 0.3.1 & jaxlib from 0.1.71 to 0.3.0

2022-02-11

- Improve keyboard navigation for the open dialog
- Fix issue where nvidia-smi stopped reporting resource utilization for some users who were modifying the version of nvidia used
- Update tensorflow from 2.7 to 2.8, keras from 2.7 to 2.8, numpy from 1.19.5 to 1.21.5, tables

Colab notebooks and notebooks opened from GitHub Gists
- Update pandas from 1.1.5 to 1.3.5
- Update openpyxl from 2.5.9 to 3.0.0 and pyarrow from 3.0.0 to 6.0.0
- Link to the release notes from the Help menu

2022-01-28

- Add a copy button to data tables
- Python LSP support for better completions and code diagnostics. This can be configured in the Editor Settings (Tools → Settings)
- Update gspread examples in our documentation
- Update gdown from 3.6 to 4.2

2022-01-21

- New documentation for the `google.colab` package
- Show GPU RAM in the resource usage tab
- Improved security for mounting Google Drive which disallows mounting Drive from accounts other than the one currently executing the notebook

2022-01-14

- Add a preference (Tools → Settings) to use a temporary scratch notebook as the default landing page
- Fix bug where / and : weren't working in VIM